

VIAP - A System for Verifying Integer Assignment Programs with Loops

Pritom Rajkhowa and Fangzhen Lin

Department of Computer Science
The Hong Kong University of Science and Technology
Clear Water Bay, Kowloon, Hong Kong,
`{prajkhowa,flin}@cse.ust.hk`

Abstract. This paper describes a simple implemented system for proving the correctness of a procedural program with integer assignments and loops. It translates a given program to a set of first-order axioms with quantification over natural numbers, and passes the resulting axioms, together with the given precondition and postcondition, to the SMT solver Z3. It successfully proves benchmark programs such as those for computing integer square roots, integer product, integer power, factorial and the sums of series of integers. The system is fully automatic and points to a new way of proving properties of programs.

Keywords: Automatic Program Verification, Preconditions, Postconditions, First-Order Logic, Mathematical Induction, Recurrences, SMT, Arithmetic

1 Introduction

Program verification has been a key problem in computer science since its beginning. Various approaches have been proposed for various types of programs. This paper considers a basic class of imperative programs composed of integer assignments, sequence, conditionals and loops. It describes the theory and implementation of a fully automatic system we developed for proving the correctness of these programs with respect to a specification given by a pair of precondition and postcondition.

We call our system VIAP, short for a Verifier for Integer Assignment Programs. Given a program without any annotations, and a proof obligation in the form of a precondition and a postcondition, VIAP first translates the program to a set of first-order axioms with natural number quantifications. This part is independent of the proof obligation. It then attempts to prove the postcondition from the translated axioms and the precondition using the SMT solver Z3 [1] as the basic theorem prover.

The translation part is based on a recent proposal by Lin [2]. We improved it by using off-the-shelf systems sympy [3] and Mathematica [4] to try to compute closed-form solutions of inductive definitions in Lin's translation. The theorem proving part uses Z3 directly when the translated axioms have no inductive

definitions. Otherwise, it tries a simple inductive proof using Z3 as the base theorem prover.

To illustrate how our system works, consider the following program for computing the integer square root of a given non-negative integer X :

```
a=0; su=1; t=1;
while ( su <= X ) {
  a=a+1; t=t+2; su=su+t; }
return a;
```

With some simple simplification, the translation outlined in [2] would generate the following axioms:

$$\begin{aligned}
X_1 &= X, \\
a_1 &= a_3(N), \quad a_3(0) = 0, \quad a_3(n+1) = a_3(n) + 1, \\
t_1 &= t_3(N), \quad t_3(0) = 1, \quad t_3(n+1) = t_3(n) + 2, \\
su_1 &= su_3(N), \quad su_3(0) = 1, \quad su_3(n+1) = su_3(n) + t_3(n) + 2, \\
&\neg(su_3(N) \leq X), \\
&(n < N) \rightarrow (su_3(n) \leq X),
\end{aligned}$$

where a denotes the input value of the program variable \mathbf{a} , a_1 is the output value, and $a_3(n)$ a temporary function denoting the value of \mathbf{a} after the n th iteration of the while loop. Similar conventions apply to X , t , su and their subscripted versions. The constant N is introduced to denote the number of iterations of the while loop before it exits. The variable n ranges over natural numbers and is universally quantified.

By using sympy and if necessary, Mathematica, our translator computes closed-form solutions for $a_3()$, $t_3()$, and $su_3()$, eliminates them, and produces the following axioms:

$$\begin{aligned}
X_1 &= X, \quad a_1 = N, \quad t_1 = 2N + 1, \quad su_1 = N^2 + 2N + 1, \\
&\neg(N^2 + 2N + 1 \leq X), \\
&(n < N) \rightarrow (n^2 + 2n + 1 \leq X).
\end{aligned}$$

With these set of axioms, Z3 can then be made to prove the following two postconditions

$$(a_1 + 1)^2 > X, \quad a_1^2 \leq X$$

given the precondition $X \geq 0$.

This is an example where all recurrences happen to have closed-form solutions that can be computed by both Mathematica and sympy. Sometimes there are recurrences that can be solved by Mathematica but not by sympy and vice versa. More often, they cannot be solved by either of them. When this happens, the proof cannot be proved directly by Z3 as it needs to use mathematical induction which is not inbuilt in Z3. There has been work on extending Z3 with induction

[5], and more recently extending CVC4 with induction [6]. For now our system implements a simple induction scheme.

Our main motivation for constructing VIAP is to show that the translation proposed in [2] can work using off-the-shelf systems such as sympy and Z3. Specifically, VIAP implements and extends the translation proposed in [2] by using sympy and Mathematica to solve recurrences generated by the translation. It then proves the partial correctness of a program with respect to a given precondition and post-condition by mathematical induction with Z3 as the base prover. We have used it to prove the correctness of many benchmark programs from [7] [8] [9] [10] [11] - the details are given in the experiment section later in the paper.

We will discuss related work later in a separate section. For now, we mention that for these benchmark programs, our system works without explicit loop invariants. Many of the programs that we tried with our systems are benchmark programs for loop invariant discovery systems. Now our systems can prove the (partial) correctness of these programs without needing these loop invariants.

The rest of this paper is organized as follows. We next describe how our system translates programs to first-order logic in more details. We then describe how we use Z3 for program verification using the translated axioms, and define two proof strategies. We next provide the experimental results of running our system on 28 programs, many of them from work by others, and some classic arithmetic algorithms, such as factorial, power, and product are included to illustrate certain features of our system. We also ran Dafny [12], Leon [10], 2LS[13], Automizer[14] and Kojak[15] on the same programs and will discuss the results. We then describe possible uses of our system for proving properties about program termination, discuss more related work and finally conclude the paper.

2 Translation

Our system handles programs consisting of integer assignments, sequence, conditionals and while loops:

```
<P> ::= <X> = <E> |
      <P>; <P> |
      if <C> then <P> else <P> |
      while <C> {<P>}
```

where $\langle X \rangle$ is an integer variable, $\langle E \rangle$ a integer expression and $\langle C \rangle$ a boolean condition.

Our translation is based on the one described in [2]. It translates a program to a set of first-order axioms by introducing for each loop a natural number term that denotes the number of iterations the loop ran before it exits. For a while loop like `while C { B }`, it generates the following two formulas:

$$\neg C(N), \tag{1}$$

$$\forall n. n < N \rightarrow C(n) \tag{2}$$

where N is a new natural number constant denoting the number of iterations that the loop ran before it exits, and $C(n)$ a formula denoting the true value of the condition C at the end of the n th iteration of the loop.

We will not give the details of the translation here but illustrate it by examples and show how we improve on it. Consider the following algorithm for computing the integer division by Cohen [16] (the one below is from [2] which is adapted from [17]). It has two loops, one nested inside another:

```
// X and Y are two input integers; Y > 0
Q=0; R=X;
while (R >= Y) do {
  A=1; B=Y;
  while (R >= 2*B) do {
    A = 2*A;
    B = 2*B;
  }
  R = R-B;
  Q = Q+A
} // return Q = X/Y;
```

Again for a variable V , use V to denote its input value, V_1 its output value, and V_2, V_3, \dots its temporary value during the program execution, our translator outputs the following set of axioms:

$$Y_1 = Y, \quad (3)$$

$$X_1 = X, \quad (4)$$

$$A_1 = A_7(N_2), \quad (5)$$

$$q_1 = q_7(N_2), \quad (6)$$

$$r_1 = r_7(N_2), \quad (7)$$

$$B_1 = B_7(N_2), \quad (8)$$

$$r_7(n) < 2 \times 2^{N_1(n)} \times Y, \quad (9)$$

$$n_1 < N_1(n_2) \rightarrow r_7(n_2) \geq 2 \times 2^{n_1} \times Y, \quad (10)$$

$$A_7(n+1) = 2^{N_1(n)} \times 1, \quad (11)$$

$$B_7(n+1) = 2^{N_1(n)} \times Y, \quad (12)$$

$$q_7(n+1) = q_7(n) + 2^{N_1(n)} \times 1, \quad (13)$$

$$r_7(n+1) = r_7(n) - 2^{N_1(n)} \times Y, \quad (14)$$

$$A_7(0) = A, \quad (15)$$

$$B_7(0) = B, \quad (16)$$

$$q_7(0) = 0, \quad (17)$$

$$r_7(0) = X, \quad (18)$$

$$r_7(N_2) < Y, \quad (19)$$

$$n < N_2 \rightarrow r_7(n) \geq Y \quad (20)$$

where

- N_2 is a system generated natural number constant denoting the number of iterations that the outer loop runs before exiting;
- N_1 is a system generated natural number function denoting that for each k , $N_1(k)$ is the number of iterations that the inner loop runs during the k th iteration of the outer loop;
- n , n_1 and n_2 are universally quantified natural number variables.
- r_7 and others are system generated temporary constants and functions denoting the values of the corresponding program variables during program execution. For this program, for each variable, say r , our translator generates six temporary ones r_2, \dots, r_7 and eliminate all of them except the last one.

We use sympy to simplify algebraic equations and solve recurrences, and use Mathematica as a backup for solving recurrences when sympy could not. For this program, the inner loop generates two recurrences, one for A and the other for B :

$$\begin{aligned} A_2(n+1) &= 2 \times A_2(n), & A_2(0) &= A, \\ B_2(n+1) &= 2 \times B_2(n), & B_2(0) &= B. \end{aligned}$$

sympy solves both, and they are replaced by their closed form solutions:

$$A_2(n) = 2^n \times A, \quad B_2(n) = 2^n \times B.$$

This is the reason why we have (11) and (12).

3 Proof strategies

We use Z3 [1] to prove properties about a program with the axioms generated from our translator. Given that Z3 accepts first-order axioms, this seems like a straightforward task. The problem is of course that theorem proving in first-order logic is undecidable, and there are many limitations in what Z3 can prove. We describe below the strategies that we use to make Z3 work for us.

First, we use a natural number sort, but Z3 has only integer sort. So whenever we need to quantify over a natural number variable, we need to make it into a quantification on non-negative integers. For example, to encode in Z3 the following axiom that universally quantifies natural number variable n :

$$n < N_2 \rightarrow r_7(n) \geq Y$$

we use an integer type variable n and add the condition $n \geq 0$:

```
n=Int('n')
ForAll([n], Implies(n>=0, Implies(n<N2,r7(n)>=Y)))
```

Another issue is that many of the benchmark programs that we tried have exponentiation. We found that Z3's builtin module NLSat (nonlinear real arithmetic) was too weak most of the time. Instead, we define our own exponentiation function $power(x, y)$ on integers:

```
power=Function('power',IntSort(),IntSort(),IntSort())
```

and include the following axioms:

```
ForAll([n],Implies(n>=0, power(0,n)==0))
ForAll([n,m],Implies(power(m,n)==0,m==0))
ForAll([n],Implies(n>0, power(n,0)==1))
ForAll([n,m],Implies(power(n,m)==1,Or(n==1,m==0)))
ForAll([n,m],Implies(And(n>0,m>=0), power(n,m+1)==power(n,m)*n))
```

Of course, these axioms are not complete. But so far they work better for us than Z3's build-in module.

Yes another issue that we had with Z3 was that it has difficulty with axioms of the form (2), which is crucial for our axiomatization of loops. Our strategy is to add the following consequence of (2) to the axiom set:

$$N = 0 \vee C(N - 1). \quad (21)$$

There are some other preprocessings that we have done while translating the axioms to Z3. Chiefly, we make use of sympy's expand and simplify functions to convert arithmetic terms to a standard form. For example, a term like $(x+1)^2 + x$ is converted to $x^2 + 3x + 1$, and exponentiation are simplified by grouping those with the same base, for example, by converting $x^y \times x^z$ to x^{y+z} .

Finally, since Z3 does not have a build-in mechanism for induction, we implemented a simple induction scheme on top of it. In the end, we came up with two proof strategies: direct proof and proof by simple induction.

3.1 Strategy 1: direct proof

Input A - a set of first-order axioms; α - precondition (a sentence on the input values of program variables); β - postcondition (a sentence on the input and output values of program variables);

Output “proved”, “failed to prove”, or “refuted” (with a counter-example);

1. Declare all variables in A , α and β as integer variables;
2. Add binary integer function **power**(x, y) and its associate axioms;
3. Convert α and each axiom in A to Z3 format (after some pre-processing by sympy);
4. For each axiom in A of the form (2), add (21);
5. For each output variable $X1$ in β , match it with an equation of the form $X1 = W$ in A , and replace it in β by W . Let the resulting sentence be β' . Convert $\neg\beta'$ to Z3 format.
6. Query Z3 with the axioms thus generated and output “proved” if it returns “unsat”, “refuted” (with a counter-example) if it returns the counter-example, and “failed to prove” otherwise.

3.2 Strategy 2: simple induction

The first four steps of this strategy are the same as the first one:

Input A - a set of first-order axioms; α - precondition (a sentence on the input values of program variables); β - postcondition (a sentence on the input and output values of program variables);

Output “proved”, “failed to prove”, or “refuted” (with a counter-example);

1. Declare all variables in A , α and β as integer variables;
2. Add binary integer function **power**(x, y) and its associate axioms;
3. Convert α and each axiom in A to Z3 format (after some pre-processing by sympy);
4. For each axiom in A of the form (2), add (21);
5. Let Q be the set of Z3 axioms generated so far.
6. For each output variable $X1$ in β , match it with an equation of the form $X1 = W$ in A , and replace it in β by W . Let the resulting sentence be β' . Find in β' a system generated natural number constant N denoting the number of iterations in a loop. The following steps try to prove $\forall k. \beta'(N/n)$ by induction on n .
7. Replace N in β' by 0. Let β'' be the resulting sentence. Convert $\neg\beta''$ to Z3 axioms and add it to Q .
8. Query Z3 with the axioms thus generated. If it returns “unsat”, then continue. Otherwise, return “failed to prove”.
9. Choose a new symbol k not occurring in Q or β' , and declare it to be a non-negative integer.
10. Replace N in β' by k to get a new sentence β_1 . Replace N in β' by $k + 1$ to get another sentence β_2 .
11. For each term in β_2 of the form $f(k + 1)$, match it with an equation in A of the form $\forall n. f(n + 1) = W(n)$, and replace $f(k + 1)$ in β_2 by $W(k)$. Let the resulting sentence be β_3 . Convert $\neg(\beta_1 \rightarrow \beta_3)$ to Z3 format and add it to Q .
12. Query Z3 with the axioms thus generated. If it returns “unsat”, then return with “proved”. Otherwise, return “failed to prove”.

As an illustration, consider proving the following postcondition β (recall that q_1 and r_1 denote the output values of q and r , respectively):

$$X = q_1 Y + r_1 \quad (22)$$

under the precondition α :

$$X \geq 0 \wedge Y > 0$$

for the above Cohen’s division program. First VIAP converts the postcondition into the following assertion using the equations about q_1 and p_1 in the set of axioms:

$$X = q_7(N_2)Y + r_7(N_2).$$

Then it tries to prove it directly using Z3, but fails. So it tries to prove the following more general one

$$\forall n. X = q_7(n)Y + r_7(n) \quad (23)$$

by induction on n , which succeeds.

4 Experiments

4.1 Verifiability

VIAP is implemented in python. The source code and the full experiments are available on

<https://github.com/VerifierIntegerAssignment/VIAP>

VIAP sometimes call Wolfram Alpha to solve a recurrence. When this happens, a working internet connection is needed. Table 1 is a summary of our experiments. We discuss them in more details below.

4.2 Benchmarks

We considered 29 programs, most of them are from or adapted from E. Rodriguez-Carbonell and Kapur’s NLA test suite for their work [7] [8] [9]. In particular, the 10 power series programs `potSumm1` - `potSumm10` are adapted from the NLA test suite. These are programs for computing the power sum $\sum_{n=1}^X n^i$ for $i = 1, \dots, 10$. As shown in Table 1, for these programs, the postconditions are the closed-form formulas of the corresponding power sum. For example, for `potSumm2`, the postcondition is that the final value of `sum` will be $X*(X+1)*(2*X+1)/6$, which is a closed-form formula for $\sum_{n=1}^X n^2$. These programs demonstrate the capability of VIAP’s recurrence solving module to find closed form solution of the polynomial recurrences of various orders. They also demonstrate the effectiveness of the system to prove the postconditions involving polynomial of various orders.

Two of the programs (`divDafny` and `addDafny`) are from Dafny presented at VSTTE 2008 [11], and one (`abs`) from Leon [10].

The factorial program (`fact`) is used to demonstrate that our system can handle pre-defined functions, including finding closed-form solutions using the pre-defined functions in sympy and mathematica. Tower of Hanoi program (`hanoi`) and the exponentiation program (`power`) are used to demonstrate how our system can handle postcondition involving polynomial with arbitrary exponents.

4.3 Experimental results

Our experiments were done on a PC and 4GB of RAM running Windows 7. The system was run on Python 2.7.11 with sympy 1.0, Wolfram Alpha 2.4, and Z3 4.4.1. Both sympy and Z3 are installed locally. However, Wolfram Alpha is a service from Wolfram Research and called remotely. Not all programs need Wolfram Alpha. For those that do need it, they can only be translated if there is a working internet connection to the server of Wolfram Alpha.

For Z3, we set its timeout parameter to 60 seconds. We included two sets of times in Table 1: the T_1 column are the times for VIAP to translate the given program to axioms, and T_2 the times for it to return the proofs for the given verification tasks. Both are in CPU seconds. Notice that t_1 includes the times for

calling sympy and/or Mathematica to solve recurrences and is independent of the verification tasks. The translation needs to be done only once. On the other hand, t_2 is specific to the given verification task in terms of a precondition and a postcondition. It includes time to process the two conditions as well as the Z3 query time.

We have also included some tasks that VIAP is not able to prove currently. For example, it was not able to prove the postcondition $r1 < Y$ (the remainder is smaller than the divisor) for the program `divHard` (an integer division program). Given the translated axioms for the program, VIAP converts the postcondition to $r7_N2 < Y$. The direct proof strategy could not prove it as Z3 was not powerful enough for it. Our simple induction strategy would try to prove it by proving that for all natural number n , $r7(n) < Y$ holds. But this cannot be proved as even the case, $r7(0) < Y$ may not be true.

4.4 A comparison with a few other systems

As an illustration of how VIAP compares with other program verification systems, we ran Leon [10], Dafny [12], 2LS [13], and Ultimate Kojak [15] on the same set of programs. The last four columns of Table 1 summarize the results of these systems. The full experimental results can be found in https://github.com/VerifierIntegerAssignment/VIAP/tree/master/Experimental_Data. We briefly describe these systems below.

Leon started as a verification system for the pure functional subset of Scala. Later it was extended to include imperative programs by translating them internally into pure functional programs. In our experiments, we use Leon 3.0, the latest version available on <http://lara.epfl.ch/w/leon> at the time of writing. Leon also uses Z3.

Dafny is a special language designed for formal verification. It is build directly on top of Z3. Given a program with proof obligations, it first translates it to a Boogie program which is then converted to some verification conditions by Boogie program verifier and passed to Z3. In our experiment, we use Dafny 1.9.7.30401, the latest version available on <https://dafny.codeplex.com/> at the time of writing.

2LS is a static analysis and verification tool for C program which got the highest score in the Loop subcategory of the Integers and Control Flow category in the SV-COMP'2016 competition. It is built upon the CPROVER infrastructure which implements incremental bounded model checking, invariant generation techniques and k-induction. In our experiment, we use 2LS-0.3 as submitted to SV-COMP'2016 <http://www.cprover.org/svn/deltacheck/releases/2ls-0.3-sv-comp-2016/>. 2LS uses MiniSAT 2.2.1 [18].

Ultimate Kojak is a symbolic model checker for program C program which prove the correctness of the program by computing inductive invariant from interpolation. It also performed well in the SV-COMP'2016 competition and uses Z3. We used its latest version (v0.1.6) downloaded on <https://monteverdi.informatik.uni-freiburg.de/tomcat/Website/?ui=tool&tool=kojak>. We also tried Ultimate Atomizer from the same research group with the same results.

Table 1. Experimental Results

Program	Desc	Postcondition	VIAP		T_1	T_2	T_3	T_4
			t_1 (s)	Strategy				
abs	absolute value [10]	$a1 \geq 0$	< 1	✓ (Direct)	< 1	?	×	×
sqrt	square root [7] [8] [9]	$(a1+1)**2 > X$ $a1**2 \leq X$	1	✓ (Direct) ✓ (Direct)	1	×	×	×
fact	factorial	$F1 == \text{factorial}(X)$	9	✓ (Direct)	< 1	×	×	×
product	multiply two numbers	$A1 == a * b$	< 1	✓ (Direct)	< 1	×	×	×
power	exponential	$P1 == a ** b$	< 1	✓ (Direct)	1	×	×	×
square	square of a number	$s1 == M ** 2$	< 1	✓ (Direct)	2	×	×	×
cubeCohen	cube [7] [8] [9]	$m1 == X ** 3$	6	✓ (Direct)	1	×	×	×
hanoi	Tower of Hanoi	$h1 == 2 ** x - 1$	< 1	✓ (Direct)	1	×	×	×
subDafny	subtract two numbers	$r1 == X - Y$	7	✓ (Direct)	< 1	×	×	×
addDafny	add two numbers [11]	$r1 == X + Y$	7	✓ (Direct)	< 1	×	×	×
potSumm1	geo series [7] [8] [9]	$\text{sum1} == X * (X+1) / 2$	< 1	✓ (Direct)	10	×	×	×
potSumm2	geo series [7] [8] [9]	$\text{sum1} == X * (X+1) * (2 * X + 1) / 6$	13	✓ (Direct)	15	×	×	×
potSumm3	geo series [7] [8] [9]	$\text{sum1} == X * (X+1) * X * (X+1) / 4$	16	✓ (Direct)	20	×	×	×
potSumm4	geo series [7] [8] [9]	$\text{sum1} == \alpha_4$	18	✓ (Direct)	16	×	×	×
potSumm5	geo series [7] [8] [9]	$\text{sum1} == \alpha_5$	24	✓ (Direct)	19	×	×	×
potSumm6	geo series	$\text{sum1} == \alpha_6$	27	✓ (Direct)	20	×	×	×
potSumm7	geo series	$\text{sum1} == \alpha_7$	37	✓ (Direct)	18	×	×	×
potSumm8	geo series	$\text{sum1} == \alpha_8$	46	✓ (Direct)	24	×	×	×
potSumm9	geo series	$\text{sum1} == \alpha_9$	55	✓ (Direct)	23	×	×	×
potSumm10	geo series	$\text{sum1} == \alpha_{10}$	71	✓ (Direct)	25	×	×	×
sumOfEven	geo series	$\text{sum1} == X * (X+1)$	< 1	✓ (Direct)	1	×	×	×
sumOfOdd	geo series	$\text{sum1} == X * X$	< 1	✓ (Direct)	1	×	×	×
geoReihe1	power series [7] [8] [9]	$m1 == ((Z ** K - 1) / (Z - 1)) * (Z - 1)$	1	✓ (Direct)	3	?	×	×
geoReihe2	power series [7] [8] [9]	$m1 == (Z ** K - 1) / (Z - 1)$	< 1	✓ (Direct)	1	?	×	×
geoReihe3	power series [7] [8] [9]	$m1 == ((Z ** K - 1) / (Z - 1)) * a$	< 1	✓ (Direct)	8	?	×	×
divDafny	int div [11]	$X == Y * q1 + r1$	< 1	✓ (Direct)	< 1	×	×	×
		$r1 < Y$		✓ (Direct)	< 1	✓	✓	✓
		$r1 \geq 0$		✓ (Direct)	< 1	×	×	✓
divCohen	int div [7] [8] [9]	$X == Y * q1 + r1$	2	✓ (Induction)	115	×	×	×
		$r1 < Y$		✓ (Direct)	2	✓	✓	✓
		$r1 \geq 0$		✓ (Induction)	3	?	✓	✓
divHard	int div [7] [8] [9]	$X == Y * q1 + r1$	< 1	✓ (Induction)	143	?	×	×
		$r1 < Y$		×	-	×	×	×
		$r1 \geq 0$		✓ (Induction)	127	×	×	✓
divKaldewaij	int div [7] [8] [9]	$A == B * q1 + r1$	< 1	×	-	×	×	×
		$r1 < B$		×	-	×	×	×
		$r1 \geq 0$		✓ (Induction)	120	×	×	✓

Legends:

✓: proved; ×: failed to prove; ?: Leon gave counter examples (see the text for discussion);

T_1 : Leon; T_2 : Dafny; T_3 : 2LS; T_4 : Kojak

t_1 is the time (in milliseconds) for translating the program to axioms;

t_2 is the time (in milliseconds) for proving the postcondition using the axioms;

α_4 is $(6 * X ** 5 + 15 * X ** 4 + 10 * X ** 3 - X) / 30$;

α_5 is $(2 * X ** 6 + 6 * X ** 5 + 5 * X ** 4 - X ** 2) / 12$;

α_6 is $(6 * X ** 7 + 21 * X ** 6 + 21 * X ** 5 - 7 * X ** 3 + X) / 42$;

α_7 is $(3 * X ** 8 + 12 * X ** 7 + 14 * X ** 6 - 7 * X ** 4 + 2 * X ** 2) / 24$;

α_8 is $(10 * X ** 9 + 45 * X ** 8 + 60 * X ** 7 - 42 * X ** 5 + 20 * X ** 3 - 3 * X) / 90$;

α_9 is $(2 * X ** 10 + 10 * X ** 9 + 15 * X ** 8 - 14 * X ** 6 + 10 * X ** 4 - 3 * X ** 2) / 20$;

α_{10} is $(6 * X ** 11 + 33 * X ** 10 + 55 * X ** 9 - 66 * X ** 7 + 66 * X ** 5 - 33 * X ** 3 + 5 * X) / 66$.

On some tasks, Leon returned counter-examples. These are marked by “?” in Table 1. For example, for the program `abs`:

```
object abs {
  def abs(X: Int): Int = {
    if(X <= 0) -X else X
  } ensuring(res => res >= 0)
}
```

Leon returned $X = -2147483648$ as a counter-example. The reason is that Leon assumes 32-bit representation of integers: $2^{31} = 2147483648$. Indeed, once we restrict the input X by the following condition:

```
require(-2147483648 < X && X < 2147483647)
```

Leon was able to prove the post-condition. Although 2LS and Kojak did not return counter examples, they succeeded under the restricted input in precondition like Leon. The reason is that Leon, 2LS, and Kojak assume either a 32-bit or 64-bit architecture memory model. In contrast, Dafny and our system VIAP assume an “ideal” integer model without any bounds. Thus in our model, for example, $X^2 \geq 0$ will be true for all integers, X . In fact, we inherit this model from Z3’s built in theory for arithmetic. We also implicitly assume this by making use of sympy and Mathematica.

One could argue that it is not meaningful to compare systems that assume an infinite integer model with those that assume a finite integer model. However, there may be some interesting relationships between the two models. For instance, one way of using a system like ours for a finite integer model is to add a separate step of verifying that there is no integer overflow exception during the execution. Among the programs in our experiments, this is the case for `divDafny`, `divCohen`, `divHard` and `divKaldewaij`. Thus once they are verified under an ideal integer model, they work for all finite integer representations.

5 Partial correctness and total correctness

Typically VIAP proves only partial correctness of a program with respect to a precondition and a post-condition. This is because by introducing a natural number term to denote the number of iterations that a loop ran before it exits, we assume that the loop terminates. Otherwise, a contradiction will arise.

Consider the following program:

```
while (X != 1) {
  if (X==0) { X=1; }
}
```

Clearly the program terminates only on $X = 1$ or $X = 0$. VIAP translates it into the following axioms:¹

¹ Notice that there are two equalities used in the axioms: “=” and “==”. They have the same meaning in logic.

```

X1 = X2(_N1)
X2(_n1+1) = ite(X2(_n1)==0,1,X2(_n1))
X2(0) = X
X2(_N1)==1
_n1<_N1 -> X2(_n1)!=1

```

where `ite(c,e1,e2)` is the if-then-else conditional expression: if `c` then `e1` else `e2`. This set of axioms is contradictory with any value of X except 1 or 0. Thus given a precondition like $X = 2$, any postcondition can be proved. So in this case our system proves only partial correctness.

On the other hand, this example also shows that our translated axioms can be used to deduce on which input the program terminates. First, the set of axioms is consistent. This means that the program terminates on *some* inputs. Furthermore, it entails $X = 1 \vee X = 0$. This means that it can terminate only with these two values.

Both of these can be easily verified by Z3. The consistency is verified by Z3 as it returns a model for the following Z3 query which is a direct translation of the above axioms in Z3's python format:

```

from z3 import *
X=Int('X')
X1=Int('X1')
X2=Function('X2',IntSort(),IntSort())
_N1=Const('_N1',IntSort())
_n1=Int('_n1')
_s=Solver()
_s.set("timeout",60000)
_s.add(X1 == X2(_N1))
_s.add(ForAll([_n1],Implies(_n1>=0, X2(_n1+1)==If(X2(_n1)==0,1,X2(_n1)))))
_s.add(X2(0) == X)
_s.add((X2(_N1)==1))
_s.add(ForAll([_n1], Implies(And((_n1<_N1),_n1>=0), Not(X2(_n1)==1))))
_s.add(_N1>=0)
if sat==_s.check():
    print _s.model()
elif unsat==_s.check():
    print "unsat, proved"
else:
    print "time out"

```

Now if we add the following sentence to the above query:

```

_s.add(And(Not(X==1), Not(X==0)))

```

Z3 returns with "unsat", thus proves that the axioms entails $X = 1 \vee X = 0$.

While we have not systematically explored it, this does suggest that our translated axioms can be used for proving total correctness as well. More generally, they can be used to prove properties about termination. This will be one

direction for our future work. There are many interesting questions to consider here. For instance, if the set of axioms for a program is consistent but becomes unsatisfiable when a precondition is added, does this entail that the program will not terminate for any input that satisfies the precondition? In other words, is the inconsistency only due to non-termination?

6 Related work

There has been much work on program verification, and many related systems. Given the complexity of formal program verification, many tools provide interactive proof assistants, instead of a fully automatic prover like our system. These include ACL [19], Isabella [20], Coq [21], Boyer-Moore provers [22] and PVS [23]. Most of these interactive proof assistant systems provide facilities to operate over purely functional programming languages. Many of them also support tactic-style proofs. We hope to integrate at least some of these tactics in our system in the future.

Most of the formal verification tools for non-parallel imperative computer programs relies on having a suitable loop invariant, a formula that remains unchanged throughout the execution of the loop body, to reason about the correctness of a given loop. There has been a large amount of work on automatic discovery of loop invariants. Our system can make use of loop invariants by simply adding them to the set of axioms: an invariant I for `while C { B }` corresponds to the axiom

$$\forall n. C(n) \wedge I(n) \rightarrow I(n+1).$$

However, we want to see how much we can do without using any explicit loop invariants. Furthermore, we believe that it may be more effective to use traditional mathematical induction with first-order axioms than specialized systems with loop invariants. We have shown with VIAP that this strategy works for some non-trivial programs that would require suitable loop invariants for other systems. So far we have only implemented a straightforward induction scheme using Z3 as the base theorem prover. There has been much work on automating induction (e.g. [24]). We hope to make use of this work and implement some more powerful induction strategies in future version of VIAP.

7 Concluding remarks and future work

We have described an implemented system called VIAP that can prove the partial correctness of a structured program with integer assignments with respect to a pair of a precondition and a post-condition. VIAP works by first translating the given program to a set of first-order axioms and then using the off-the-shelf SMT prover Z3 to prove that the post-condition holds under the given precondition. Integrated into both the translation part and the proving part is the use of the off-the-shelf math systems sympy and Mathematica to solve recurrences and simplify algebraic expressions. The following are some concluding remarks:

- The translation needs to be done only once for a program and the translated axioms can be used to prove multiple post-conditions. While not shown in this paper, properties about the program during its execution can be expressed and proved as well. All one needs is a label to mark the point of interest - see [2].
- For each loop in the program, the translated axioms include an explicit representation of the exact number of iterations that this loop executes before terminating. All properties about the loops are expressed in terms of this representation.
- Loop invariants are not required. Of course, they can be added to the axioms if available. Our experiments show that for many programs, VIAP can prove properties about them that would normally require proper loop invariants. In fact, most of these programs are benchmarks for testing systems that discover loop invariants.

For future work, we want to extend the system with more powerful heuristics for induction, for wide classes of programs, and for proving properties regarding termination.

Acknowledgment

We would like to thank Jianmin Ji, Jack Yao and Haodi Zhang for useful discussions. We are grateful to the developers of Z3 and sympy for making their systems available for open use, to Wolfram Research for making Wolfram Alpha an online service available to the public. We also thank Rustan M. Leino for answering our questions regarding Dafny, and Regis Blanc for Leon. All errors remain ours.

References

1. L. de Moura and N. Bjorner, “The Z3 SMT Solver.” <http://research.microsoft.com/en-us/um/redmond/projects/z3/>, 2012.
2. F. Lin, “A formalization of programs in first-order logic with a discrete linear order,” *Artificial Intelligence*, vol. 235, pp. 1 – 25, 2016.
3. SymPy Development Team, *SymPy: Python library for symbolic mathematics*, 2016.
4. Wolfram Research Inc., *Mathematica 8.0*, 2010.
5. K. R. M. Leino, “Automating induction with an SMT solver,” in *Proceedings of the 13th International Conference on Verification, Model Checking, and Abstract Interpretation, VMCAI’12*, (Berlin, Heidelberg), pp. 315–331, Springer-Verlag, 2012.
6. A. Reynolds and V. Kuncak, “Induction for SMT solvers,” in *Verification, Model Checking, and Abstract Interpretation - 16th International Conference, VMCAI 2015, Mumbai, India, January 12-14, 2015. Proceedings* (D. D’Souza, A. Lal, and K. G. Larsen, eds.), vol. 8931 of *Lecture Notes in Computer Science*, pp. 80–98, Springer, 2015.

7. T. Nguyen, D. Kapur, W. Weimer, and S. Forrest, "Using dynamic analysis to discover polynomial and array invariants," in *Proceedings of the 34th International Conference on Software Engineering, ICSE '12*, (Piscataway, NJ, USA), pp. 683–693, IEEE Press, 2012.
8. T. Nguyen, D. Kapur, W. Weimer, and S. Forrest, "DIG: A dynamic invariant generator for polynomial and array invariants," *ACM Trans. Softw. Eng. Methodol.*, vol. 23, no. 4, pp. 30:1–30:30, 2014.
9. E. Rodríguez-Carbonell and D. Kapur, "Generating all polynomial invariants in simple loops," *J. Symb. Comput.*, vol. 42, no. 4, pp. 443–476, 2007.
10. R. Blanc, V. Kuncak, E. Kneuss, and P. Suter, "An overview of the leon verification system: Verification by translation to recursive functions," in *Proceedings of the 4th Workshop on Scala, SCALA '13*, (New York, NY, USA), pp. 1:1–1:10, ACM, 2013.
11. K. R. M. Leino and R. Monahan, "Dafny meets the verification benchmarks challenge," in *Proceedings of the Third International Conference on Verified Software: Theories, Tools, Experiments, VSTTE'10*, (Berlin, Heidelberg), pp. 112–126, Springer-Verlag, 2010.
12. K. R. M. Leino, "Dafny: An automatic program verifier for functional correctness," in *Proceedings of the 16th International Conference on Logic for Programming, Artificial Intelligence, and Reasoning, LPAR'10*, (Berlin, Heidelberg), pp. 348–370, Springer-Verlag, 2010.
13. P. Schrammel and D. Kroening, "2ls for program analysis - (competition contribution)," in *TACAS 2016*, vol. 9636 of *Lecture Notes in Computer Science*, pp. 905–907, Springer, 2016.
14. M. Chechik and J. Raskin, eds., *TACAS 2016, Held as Part of the ETAPS 2016, Eindhoven, The Netherlands, April 2-8, 2016, Proceedings*, vol. 9636 of *LNCS*, Springer, 2016.
15. E. Ermiš, A. Nutz, D. Dietsch, J. Hoenicke, and A. Podelski, "Ultimate Kojak (Competition Contribution)," in *Proceedings of the 20th TACAS*, vol. 8413 of *LNCS*, pp. 421–423, Springer, 2014.
16. E. Cohen, *Programming in the 1990s - An Introduction to the Calculation of Programs*. Texts and Monographs in Computer Science, Springer, 1990.
17. T. Nguyen, D. Kapur, W. Weimer, and S. Forrest, "Using dynamic analysis to discover polynomial and array invariants," in *2012 34th International Conference on Software Engineering (ICSE)*, pp. 683–693, IEEE, 2012.
18. N. Eén and N. Sörensson, "An extensible sat-solver," in *SAT 2003*, pp. 502–518, 2003.
19. M. Kaufmann, J. S. Moore, and P. Manolios, *Computer-Aided Reasoning: An Approach*. Norwell, MA, USA: Kluwer Academic Publishers, 2000.
20. T. Nipkow, M. Wenzel, and L. C. Paulson, *Isabelle/HOL: A Proof Assistant for Higher-order Logic*. Berlin, Heidelberg: Springer-Verlag, 2002.
21. Y. Bertot and P. Castran, *Interactive Theorem Proving and Program Development: Coq'Art The Calculus of Inductive Constructions*. Springer Publishing Company, Incorporated, 1st ed., 2010.
22. R. S. Boyer, M. Kaufmann, and J. S. Moore, "The boyer-moore theorem prover and its interactive enhancement," *Computers & Mathematics with Applications*, vol. 29, no. 2, pp. 27–62, 1995.
23. S. Owre, J. Rushby, N. Shankar, *et al.*, "Pvs specification and verification system," Look at URL <http://pvs.csl.sri.com>, 2004.
24. A. Bundy, "The automation of proof by mathematical induction," in *Handbook of Automated Reasoning (in 2 volumes)* (J. A. Robinson and A. Voronkov, eds.), pp. 845–911, Elsevier and MIT Press, 2001.